

Spring Boot glibt Gas

Virtual Threads



Motivation



- Traditional model: One thread per request → simple but resource-heavy
- Limitation of platform threads: high memory & CPU cost, expensive to create
- Scalability follows Little's Law: $\text{throughput} = \text{concurrency} / \text{latency}$
- Reactive Programming: scales well but is harder to read, debug, and profile
- Virtual threads: Lightweight, efficient, massively scalable

Java

A short horizontal bar with a color gradient from blue to green.

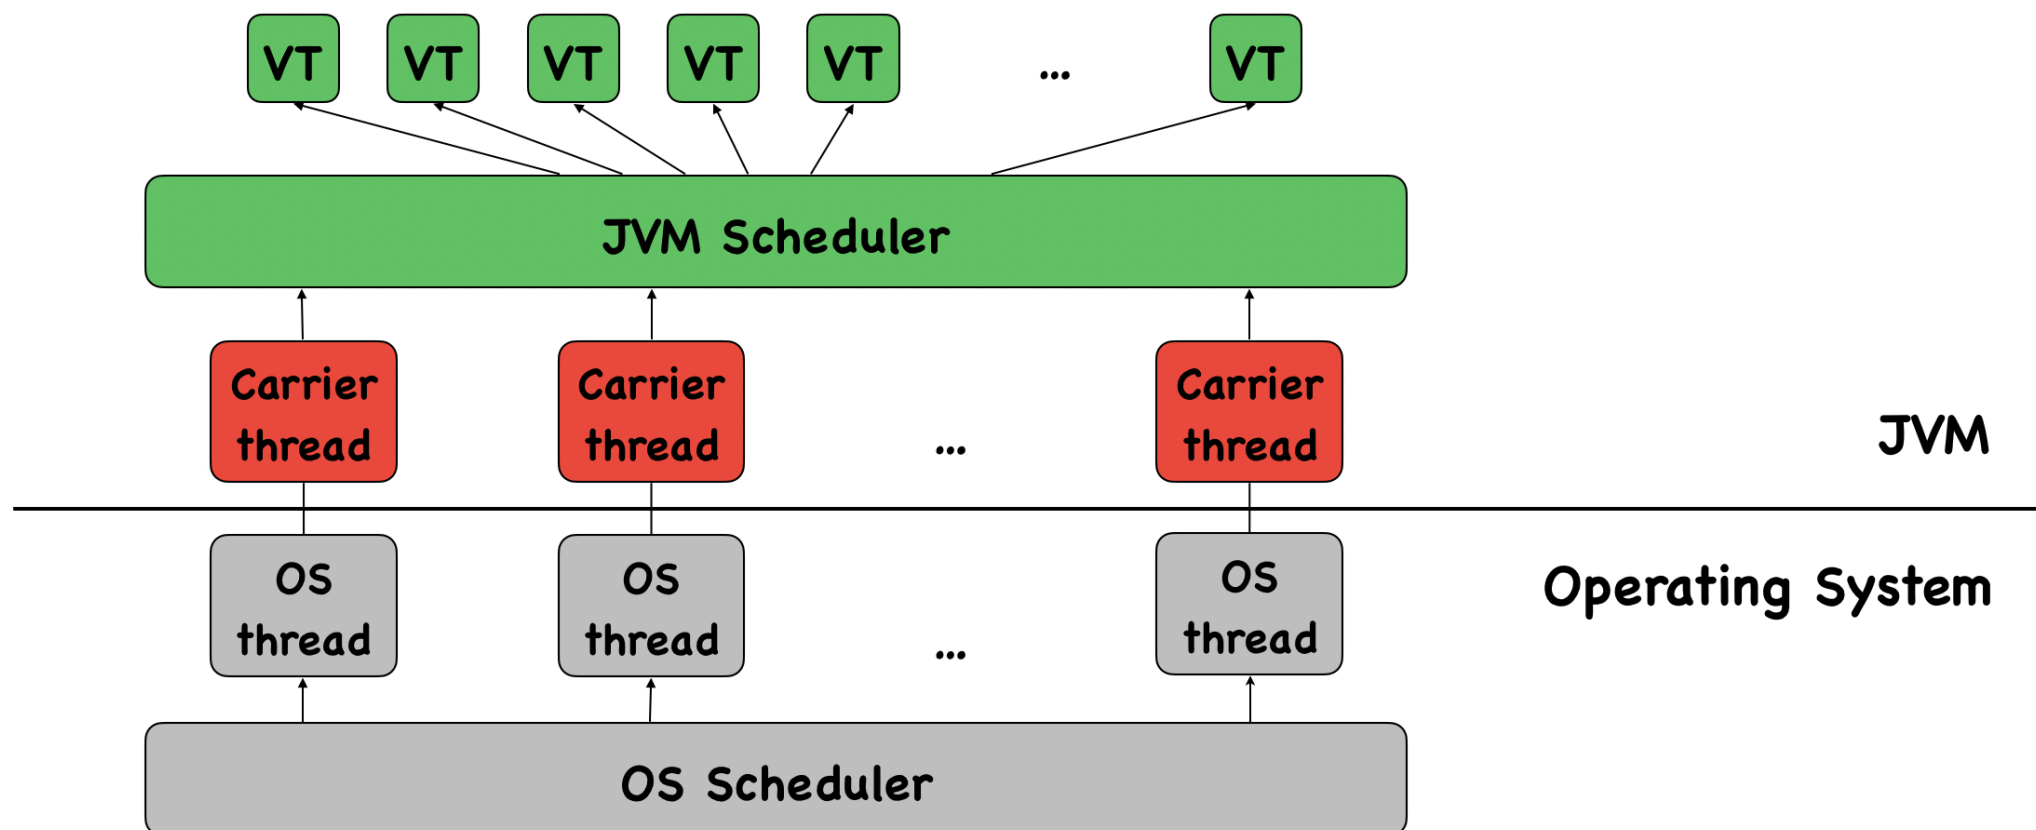
- Project Loom: introduce a lightweight concurrency construct to Java
- September '22: Virtual threads introduced as a preview feature (JDK 19, JEP 425)
- September '23: Virtual threads became stable (JDK 21, JEP 444)
- Scoped Values (JEP 446): efficient alternative to ThreadLocal for passing contextual data across threads

Spring Boot



- October '22: experimental support for virtual threads
- November '23: full support for virtual threads with release 3.2.0
- Easy to use by enabling the flag `spring.threads.virtual.enabled`
 - Tomcat and Jetty use virtual threads for request handling

Architecture



Advantages



- Ideal for I/O-bound workloads: ideal for high-throughput, I/O-heavy workloads
- Lightweight & Scalable: millions of threads with minimal memory overhead
- Simplified Concurrency Model: maintains the familiar thread-per-request model
- Better Debugging & Profiling: works with JFR, JStack, and existing tools.
- Compatibility: works with existing `java.lang.ThreadAPI`
- Improved Context Propagation: Scoped Values replace `ThreadLocal`

Disadvantages



- Not for CPU-bound Tasks: no speedup for compute-heavy workloads
- Scheduler Overhead: JVM mapping adds some cost
- Synchronization Challenges: lock contention can hurt performance
- Not a Replacement for Reactive Programming
- Requires JDK 21+ / Spring Boot 3.2.0+

Examples and Benchmark



- [Show in IDE](#)

Comparison



	Platform Threads	Virtual Threads
Implementation	Managed by the OS	Managed by the JVM (userland)
Pooling	Threads should be pooled	No pooling, new thread per task
Lifetime	Long-lived, reused	Short-lived, created per task
Usage	Multiple tasks per thread	One task per virtual thread
Thread Cost	Heavyweight (more RAM, CPU ...)	Lightweight (very low overhead)
Blocking I/O	Blocks the OS thread	Parks the virtual thread
Scalability	Limited by OS resources	Millions of concurrent threads